

MY OWN CUSTOM KAFKA CLUSTER

IN 60 SECONDS



Innovation Process Technology AG

Daniel Streb & Jonathan Gan



Welcome



whois ipt.ch

This information is subject to an Acceptable Use Policy.
See <https://www.nic.ch/terms/aup/>

Domain name:
ipt.ch

Holder of domain name:
Innovation Process Technology Inc.
Thomas Schaller
Poststrasse 14
CH-6300 Zug
Switzerland

Technical contact:
Comfox AG
Martin Fuchs

ipt AG @ ERFA



whoami #Daniel

🇨🇭 Born in Zurich

🤓 Computer Science Masters UZH

💼 Analytics, Corporate R&D, Innovation Labs

🚀 Cloud Architecture, Data, ML

❤️ Running, Hiking, Travel

whoami #Jonathan

🇸🇬 Born in Singapore

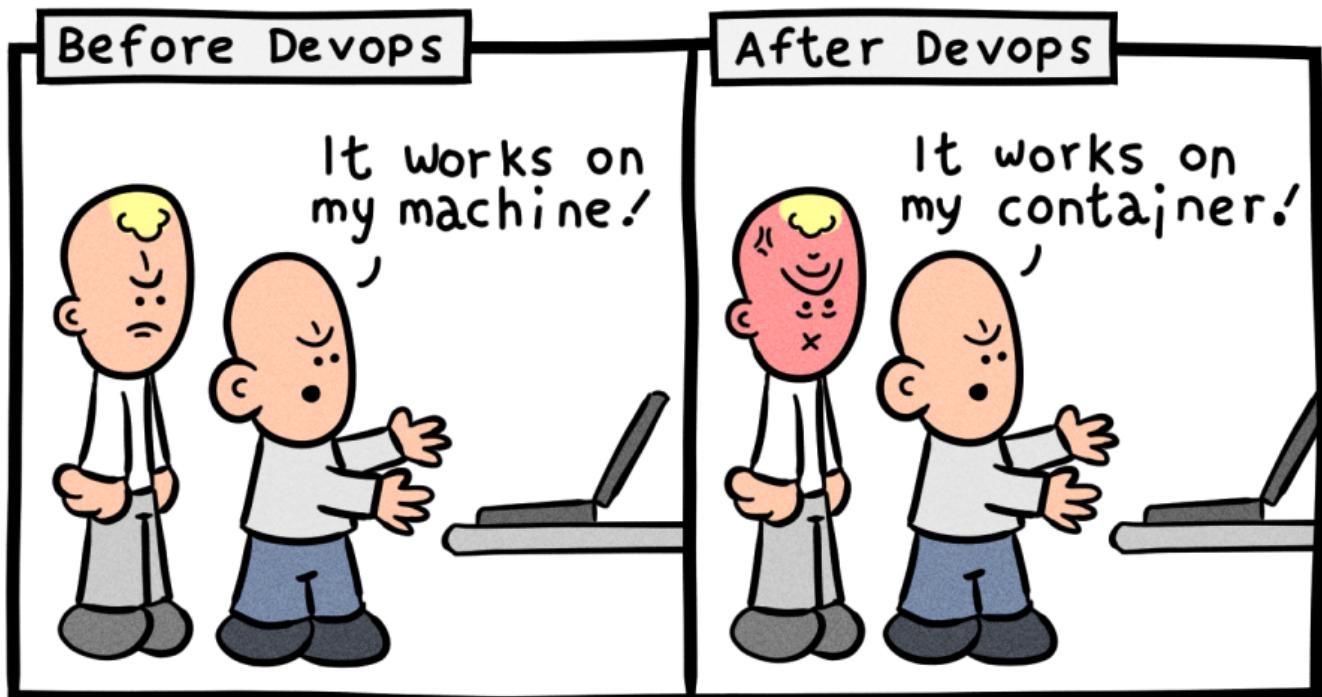
🤖 Robotics Masters ETH

💼 {Disney, IBM} Research

🚀 Computer Vision, ML, DevOps

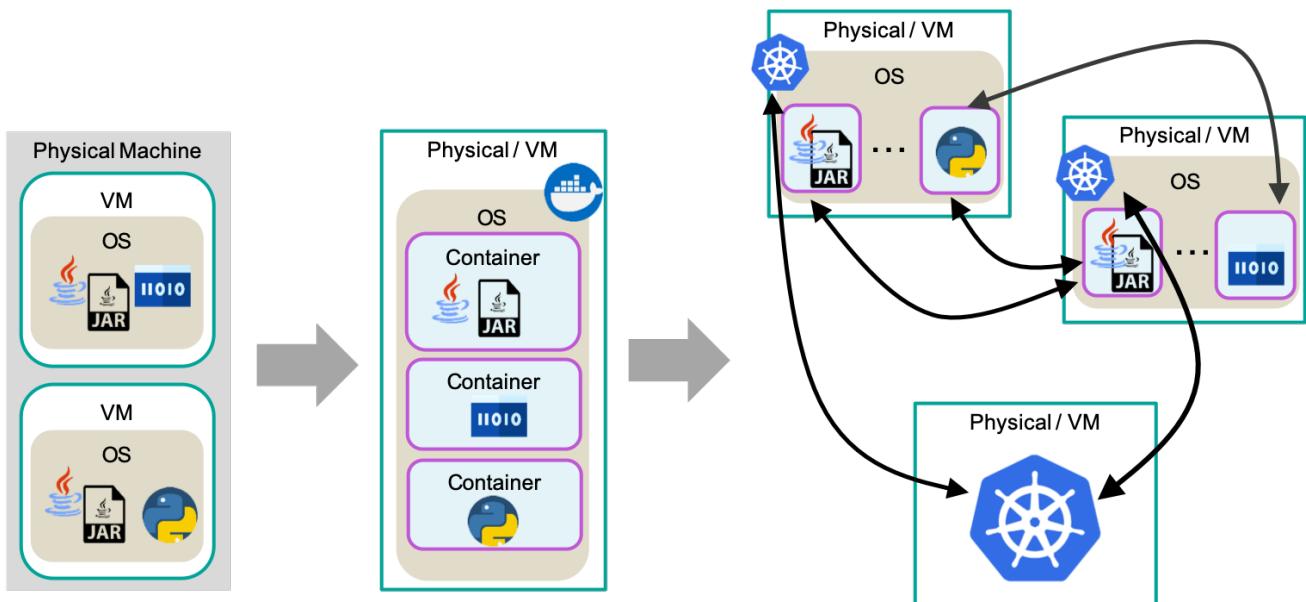
❤️ Music, Cooking, Gym

Kubernetes in a Nutshell



Daniel Stori {turnoff.us}

From Bare Metal to Container Orchestration



k8s Lingo

Pod

Container

Deployment

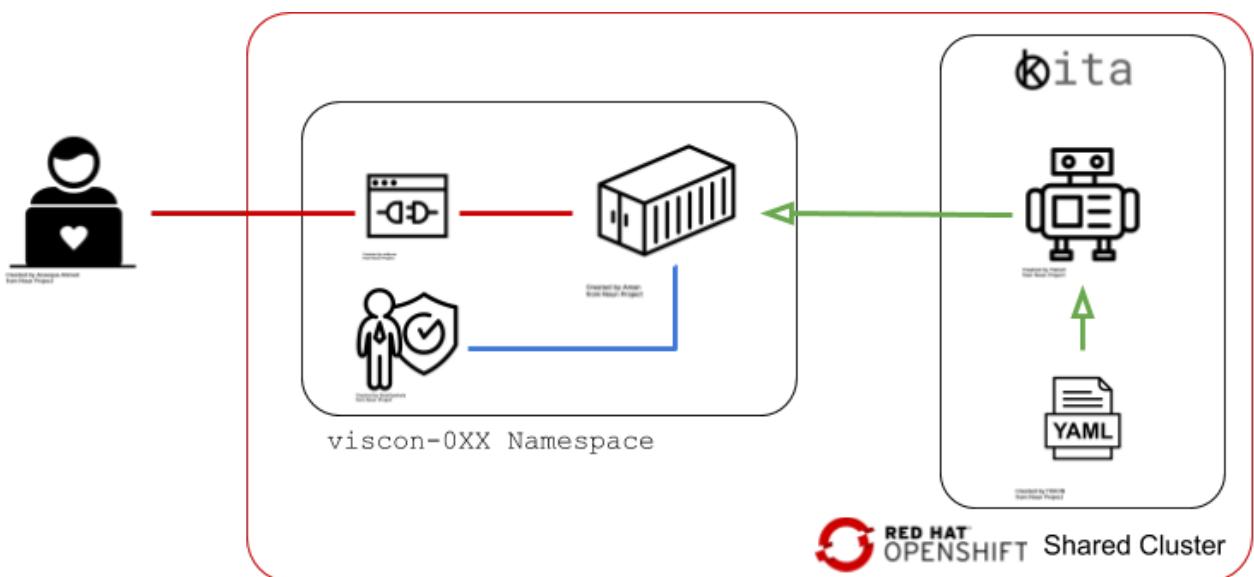
Service

Route

Serviceaccount

Intro to the Lab Environment

The high level picture



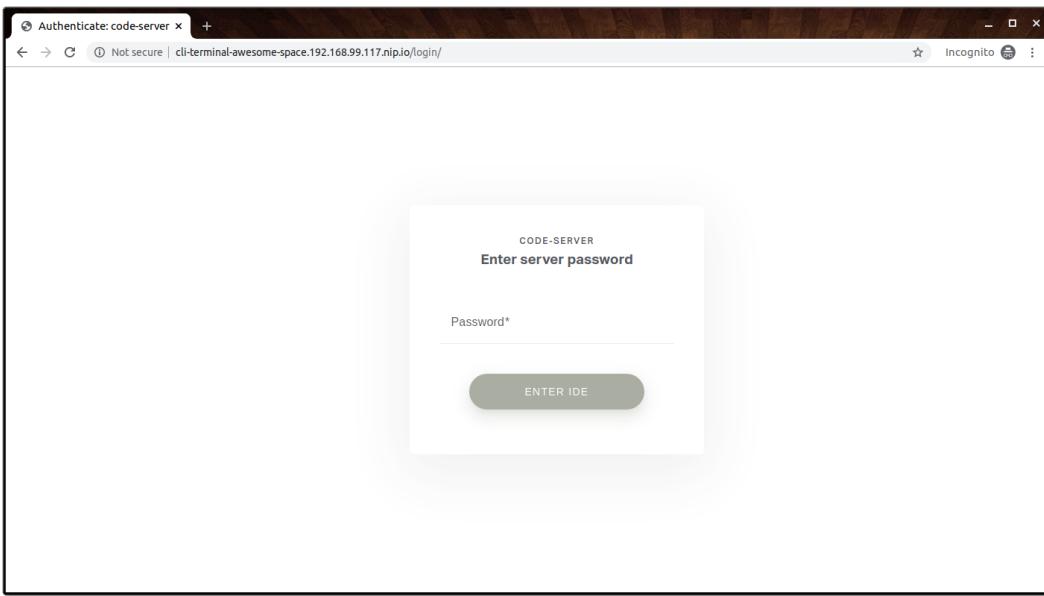
Rule 1/1: Stick to the Instructions





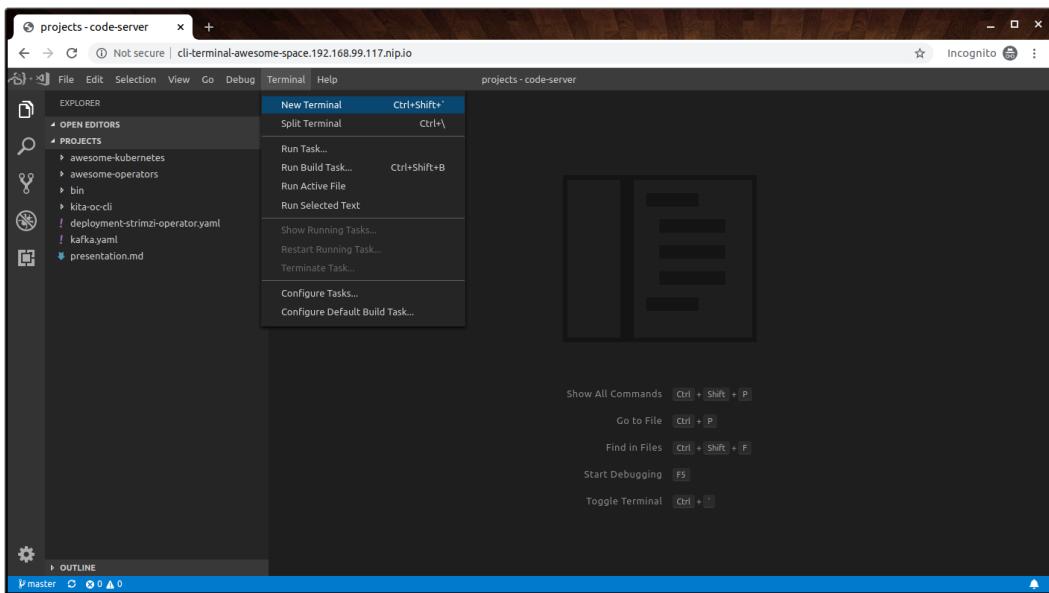
Play with your Workspace

Visit the URL on your snippet and enter the password
to access the workspace





Prepare your Shell





OpenShift CLI

Download and install the OC CLI

```
source ./kita-oc-cli/install.sh
```

Try a few things

```
oc whoami
```

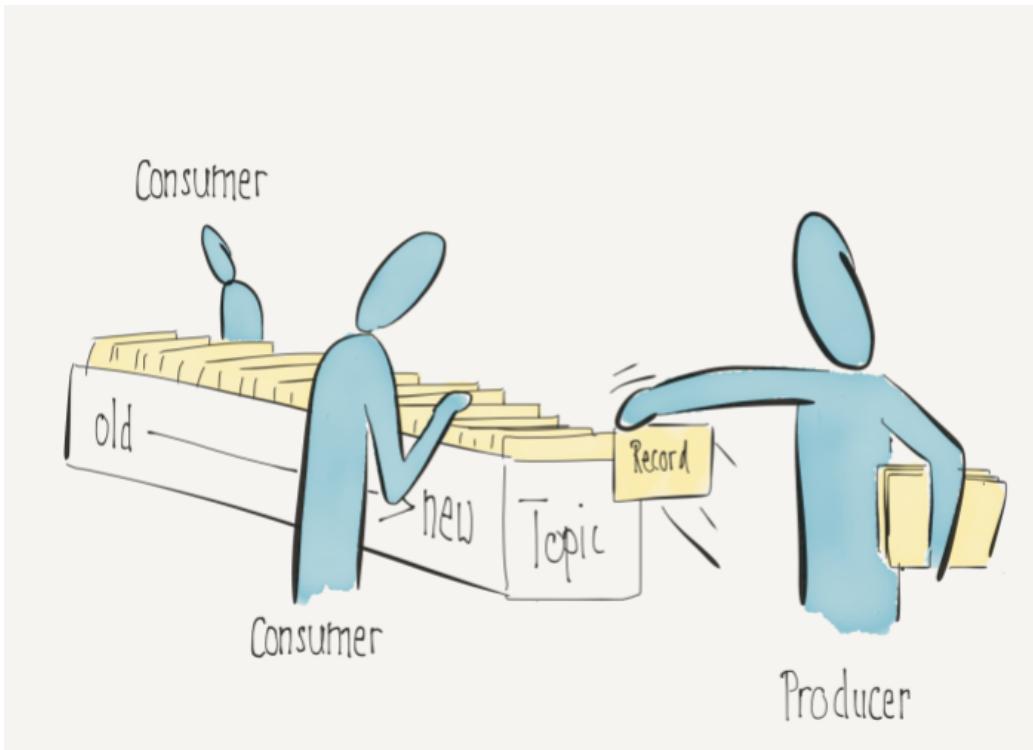
```
oc get all
```

Apache Kafka



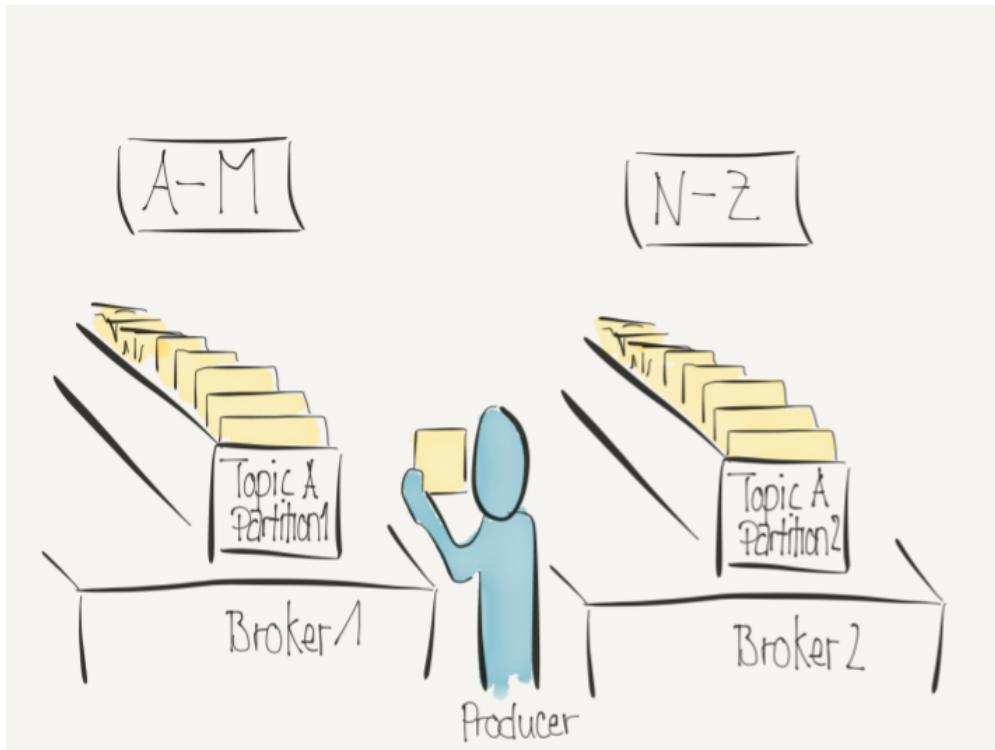
© ipt AG

Kafka High-Level



© ipt AG

Kafka Secret Sauce



© ipt AG

What is Strimzi



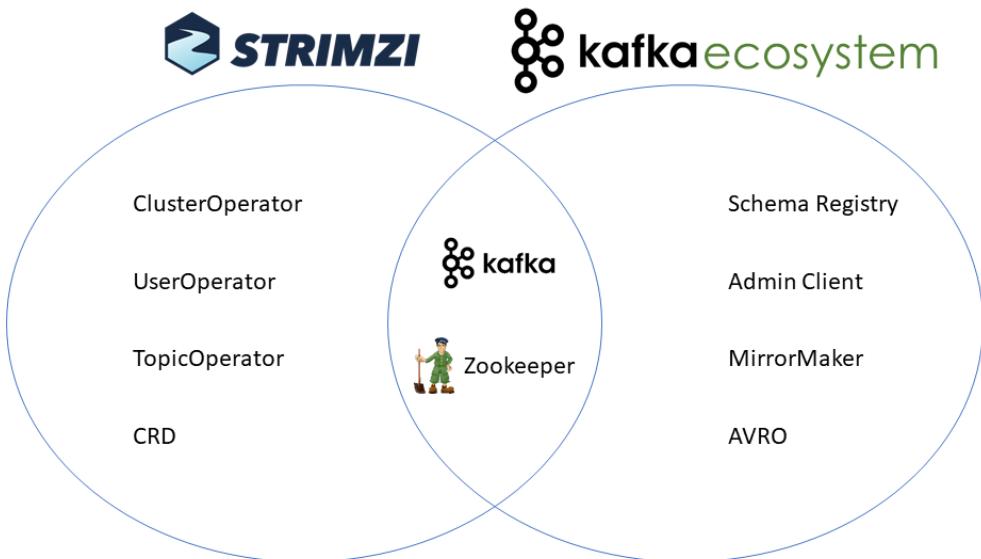
[Open Source](#) Project (started by Red Hat)

Recently adopted as a [CNCF](#) Project

Using Operators to deploy and manage Apache Kafka on Kubernetes / Openshift

Used by ipt in production since GA in 2018

Strimzi Components



Strimzi Installation

What has been done for you
(because it requires cluster admin rights)

Custom Resource Definitions

```
oc get CustomResourceDefinitions | grep kafka
```

```
kafkas.kafka.strimzi.io
kafkatopics.kafka.strimzi.io
kafkausers.kafka.strimzi.io
```

Strimzi Installation Cont.

What has been done for you
(because it requires cluster admin rights)

Roles and ClusterRoles

```
oc get ClusterRoles | grep strimzi
```

```
NAME
strimzi-cluster-operator-global
strimzi-cluster-operator-namespaced
strimzi-entity-operator
strimzi-kafka-broker
strimzi-topic-operator
```

And assigned them to the strimzi-cluster-operator service account as well as your "user"

👉 Setting up the cluster

Change into the tutorial project

```
cd /home/coder/projects/viscon-strimzi
```



Strimzi Custom Resource

```
apiVersion: kafka.strimzi.io/v1beta1
kind: Kafka
metadata:
  name: viscon-cluster
spec:
  kafka:
    version: 2.3.0
    replicas: 3
    listeners:
      plain: {}
      tls: {}
    config:
      #...
    storage:
      type: ephemeral
  zookeeper:
    replicas: 1
```

see `./strimzi/kafka/kafka-plain.yaml`

Strimzi Operator Deployment

First, we deploy the strimzi operator

```
oc apply -f strimzi/deployment-strimzi-operator.yaml
```

and wait until it is ready

```
oc get pod -l name=strimzi-cluster-operator
NAME                                     READY   STATUS    RESTARTS   AGE
strimzi-cluster-operator-6fb66c5f7-n7qr2   1/1     Running   0          10m
```

we can also inspect its logs by running

```
oc logs -f $(oc get pod -l name=strimzi-cluster-operator
```



Kafka Cluster Deployment

Create a custom Kafka resource

```
oc apply -f strimzi/kafka/kafka-plain.yaml
```

Watch the operator create a full Kafka cluster
(in about 40s 

```
oc get pods
```

NAME	READY
strimzi-cluster-operator-6bf66c5f7-n7qr2	1/1
viscon-cluster-entity-operator-6cbbd888bc-hn9fm	3/3
viscon-cluster-kafka-0	2/2
viscon-cluster-kafka-1	2/2
viscon-cluster-zookeeper-0	2/2

First End To End Example

Start simple producers and consumers to write to and read from a Kafka topic.





First Kafka Producer

Start a console producer pod

```
oc run kafka-producer -ti --image=strimzi/kafka:0.14.0-k  
--rm=true --restart=Never -- bin/kafka-console-producer  
--broker-list viscon-cluster-kafka-bootstrap:9092 \  
--topic viscon-console-demo
```

and produce a few messages

```
If you don't see a command prompt, try pressing enter.  
>Hello Viscon  
>2nd hello  
>3rd hello
```

Kill the session with CTRL + C when done



First Kafka Consumer

Start a console consumer pod

```
oc run kafka-consumer -ti --image=strimzi/kafka:0.14.0-k  
--rm=true --restart=Never -- bin/kafka-console-consumer  
--bootstrap-server viscon-cluster-kafka-bootstrap:9092  
--topic viscon-console-demo --from-beginning
```

and receive the previously produced messages

```
If you don't see a command prompt, try pressing enter.  
Hello Viscon  
2nd hello  
3rd hello
```

Kill the session with CTRL + C when done

👋 Explore Kafka Topics

Explore the Topics within Kafka

```
oc run kafka-topic-list -ti --image=strimzi/kafka:0.14.0  
--rm=true --restart=Never -- bin/kafka-topics.sh --list  
--bootstrap-server viscon-cluster-kafka-bootstrap:9092  
  
__consumer_offsets  
viscon-console-demo  
pod "kafka-topic-list" deleted
```

And the generated KafkaTopic Resources

```
oc get KafkaTopics  
  
NAME  
consumer-offsets---84e7a678d08f4bd226872e5cdd4eb527fadcl  
viscon-console-demo
```

Kafka Topic - Spec First

Take a look at the KafkaTopic in

```
strimzi/console-demo/topic.yaml
```

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaTopic
metadata:
  name: viscon-topic
  labels:
    strimzi.io/cluster: viscon-cluster
spec:
  partitions: 6
  replicas: 2
  config:
    retention.ms: 86400000 #24h
```

Now, create the topic by applying the CR

```
oc apply -f strimzi/console-demo/topic.yaml
```




Kafka Topic - Operator Magic

Let's look at the Entity Operator Logs

```
oc logs -f $(oc get pods -l strimzi.io/name=viscon-clust
```

And the Kafka internal Topic List

```
oc run kafka-topic-list -ti --image=strimzi/kafka:0.14.0  
--rm=true --restart=Never -- bin/kafka-topics.sh --list  
--bootstrap-server viscon-cluster-kafka-bootstrap:9092
```

```
__consumer_offsets  
viscon-topic  
viscon-console-demo  
pod "kafka-topic-list" deleted
```

That was cute?!

So what is missing?



Authentication

Processing

Metrics



Kafka Authentication

We no longer want our topics to be open to the world.

We can enforce [mutual tls](#)
by adapting our Kafka resource:

```
diff strimzi/kafka/kafka-plain.yaml strimzi/kafka/kafka-  
10,11c10,12  
<      plain: {}  
<      tls: {}  
---  
>      tls:  
>          authentication:  
>              type: tls
```



Kafka Authentication Cont.

To change the broker listener config we apply the

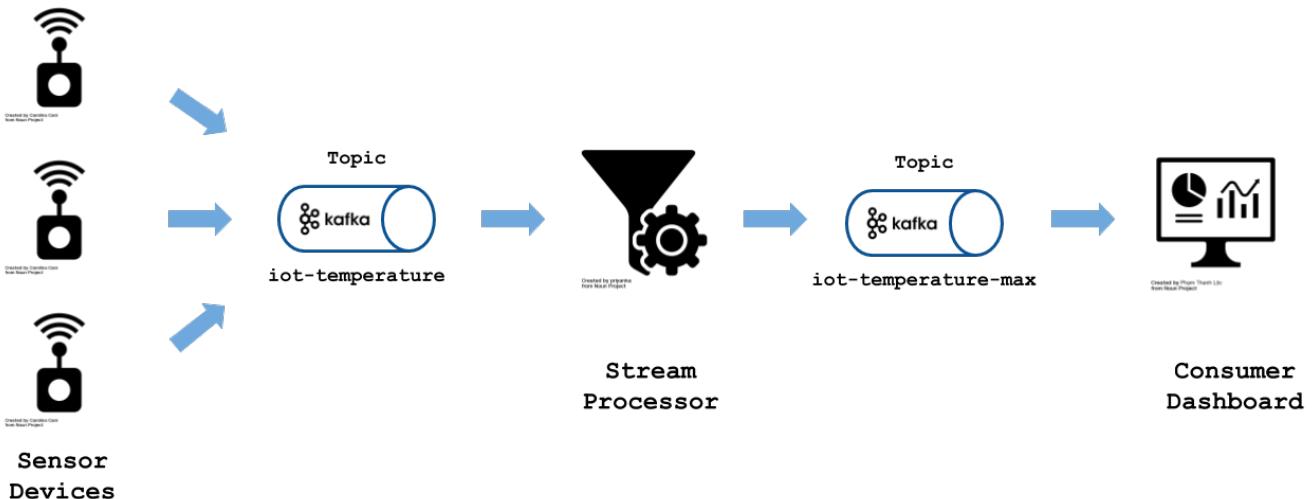
`strimzi/kafka/kafka-mtls.yaml`

config.

```
oc apply -f strimzi/kafka/kafka-mtls.yaml
```

and wait for the Kafka brokers to update

More Sophisticated Scenario





Create Demo Producer User

```
strimzi/iot-demo/iot-device/users.yaml
```

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaUser
metadata:
  name: iot-device
  labels:
    app: iot-demo
    strimzi.io/cluster: viscon-cluster
spec:
  authentication:
    type: tls
  authorization:
    type: simple
  acls:
    - resource:
        type: topic
        name: iot-temperature
        patternType: literal
```




Create Demo Stream User

```
strimzi/iot-demo/iot-stream/users.yaml
```

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaUser
metadata:
  name: iot-stream
  labels:
    app: iot-demo
    strimzi.io/cluster: viscon-cluster
spec:
  authentication:
    type: tls
  authorization:
    type: simple
  acls:
    - resource:
        type: topic
        name: iot-temperature
        patternType: literal
```




Create Demo Consumer User

```
strimzi/iot-demo/iot-consumer/users.yaml
```

```
apiVersion: kafka.strimzi.io/v1beta1
kind: KafkaUser
metadata:
  name: iot-consumer
  labels:
    app: iot-demo
    strimzi.io/cluster: viscon-cluster
spec:
  authentication:
    type: tls
  authorization:
    type: simple
  acls:
    - resource:
        type: topic
        name: iot-temperature-max
        patternType: literal
```




Create all Users

```
oc apply -f strimzi/iot-demo/iot-device/users.yaml  
oc apply -f strimzi/iot-demo/iot-stream/users.yaml  
oc apply -f strimzi/iot-demo/iot-consumer/users.yaml
```

Verify the ACL for `iot-temperature`

```
oc exec viscon-cluster-zookeeper-0 -c zookeeper -- /opt/  
--authorizer-properties zookeeper.connect=localhost:2181  
--list --topic iot-temperature
```

and for `iot-temperature-max`

```
oc exec viscon-cluster-zookeeper-0 -c zookeeper -- /opt/  
--authorizer-properties zookeeper.connect=localhost:2181  
--list --topic iot-temperature-max
```



Kafka User Secrets

The entity operator created user certificates for all users which include a CA certificate, a public certificate and a private key.

```
oc get secrets -l strimzi.io/kind=KafkaUser -o yaml
```



Mounting User Secrets

E.g. `strimzi/iot-demo/iot-device/deployment.yaml`

```
containers:
- name: device-app
  image: danistrebel/strimzi-device-app:latest
  env:
    - name: USER_CRT
      valueFrom:
        secretKeyRef:
          name: iot-device
          key: user.crt
    - name: USER_KEY
      valueFrom:
        secretKeyRef:
          name: iot-device
          key: user.key
```



Deploy the IoT Device

```
oc apply -f strimzi/iot-demo/iot-device/deployment.yaml
```

and check the logs

```
oc logs -f $(oc get pod -l name=device-app -o name | head -1)
```

Full Code

<https://github.com/danistrebel/strimzi-lab>

👉 Deploy the IoT Stream Processor

```
oc apply -f strimzi/iot-demo/iot-stream/deployment.yaml
```

and check the logs

```
oc logs -f $(oc get pod -l name=stream-app -o name | head -1)
```

Full Code

<https://github.com/danistrebel/strimzi-lab>



Deploy the Consumer

```
oc apply -f strimzi/iot-demo/iot-consumer/deployment.yaml
```

and check the logs

```
oc logs -f $(oc get pod -l name=consumer-app -o name | head -1)
```

Full Code

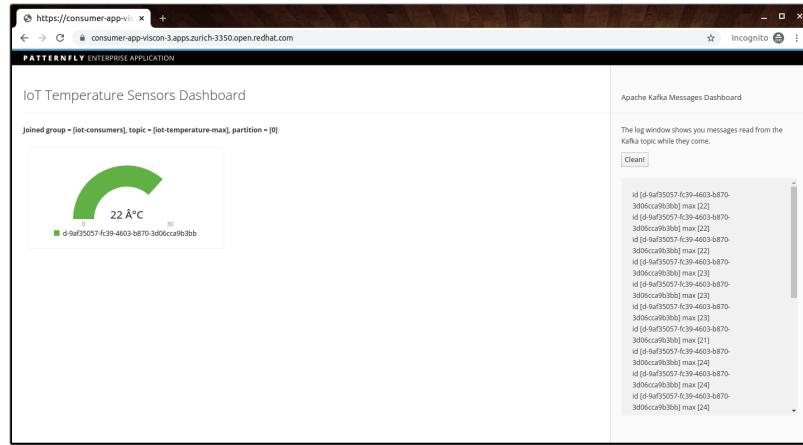
<https://github.com/danistrebel/strimzi-lab>



Check the Dashboard

```
oc get route consumer-app -o jsonpath='{.spec.host}'
```

and open it in a new Tab

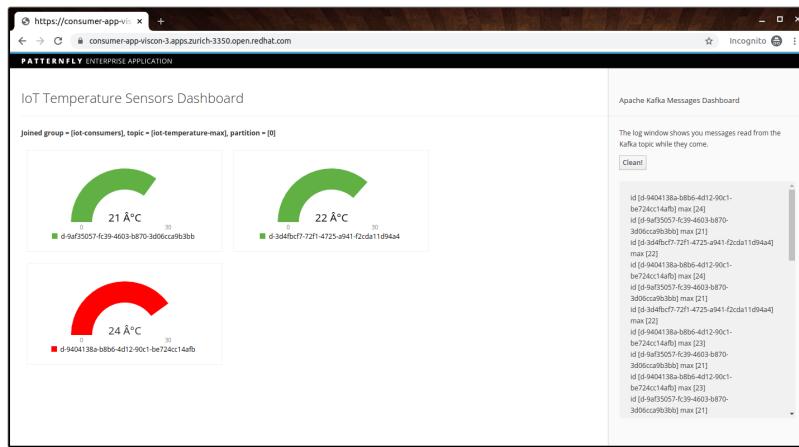


⚠️ Scale the Deployment

Increase the number of sensor devices to 3

```
oc scale deployment device-app --replicas 3
```

and wait for the data to show up in the dashboard





Cluster Monitoring

Create a scrape config

```
oc create secret generic additional-scrape-configs --from=
```

and the Prometheus deployment

```
oc apply -f strimzi/metrics-demo/prometheus.yaml
```

```
oc apply -f strimzi/metrics-demo/prometheus-service-mon
```



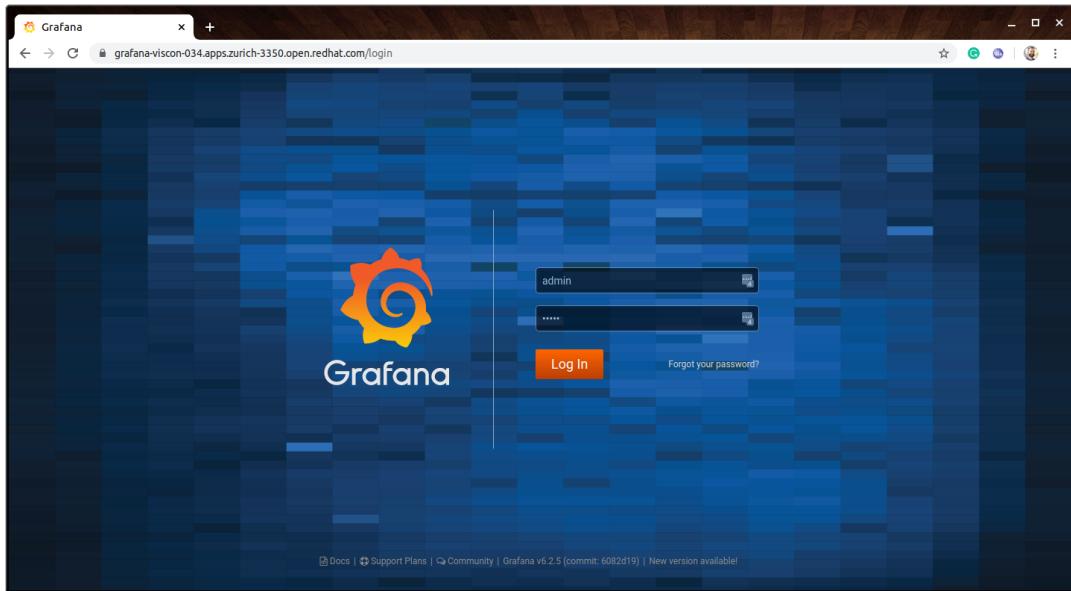
Check the Prometheus Targets

```
oc get route prometheus
```

Endpoint	State	Labels	Last Scrape	Scrape Duration	Error
https://kubernetes.default.svc:443/api/v1/nodes/infranode1.zurich-3350.internal/proxy/metrics/cadvisor	UP	beta.kubernetes.io_arch=amd64, beta.kubernetes.io_os=linux, instance=infraode1.zurich-3350.internal, job=kubernetes-cadvisor, kubernetes.io_hostname=infranode1.zurich-3350.internal, logging_infra fluentd=true, node_ip=192.168.0.229, node_name=infraode1.zurich-3350.internal, node_role=kubernetes.io_infra=true	5.922s ago	175.6ms	
https://kubernetes.default.svc:443/api/v1/nodes/master1.zurich-3350.internal/proxy/metrics/cadvisor	UP	beta.kubernetes.io_arch=amd64, beta.kubernetes.io_os=linux, instance=master1.zurich-3350.internal, job=kubernetes-cadvisor, kubernetes.io_hostname=master1.zurich-3350.internal, logging_infra fluentd=true, node_ip=192.168.0.177, node_name=master1.zurich-3350.internal, node_role=kubernetes.io_master=true	6.138s ago	135.9ms	
https://kubernetes.default.svc:443/api/v1/nodes/node1.zurich-3350.internal/proxy/metrics/cadvisor	UP	beta.kubernetes.io_arch=amd64, beta.kubernetes.io_os=linux, instance=node1.zurich-3350.internal, job=kubernetes-cadvisor, kubernetes.io_hostname=node1.zurich-3350.internal, logging_infra fluentd=true, node_ip=192.168.0.240, node_name=node1.zurich-3350.internal, node_role=kubernetes.io_comPUTE=true	8.755s ago	88.44ms	
https://kubernetes.default.svc:443/api/v1/nodes/node10.zurich-3350.internal/proxy/metrics/cadvisor	UP	beta.kubernetes.io_arch=amd64, beta.kubernetes.io_os=linux, instance=node10.zurich-3350.internal, job=kubernetes-cadvisor, kubernetes.io_hostname=node10.zurich-3350.internal, logging_infra fluentd=true, node_ip=192.168.0.241, node_name=node10.zurich-3350.internal, node_role=kubernetes.io_comPUTE=true	3.745s ago	92.94ms	

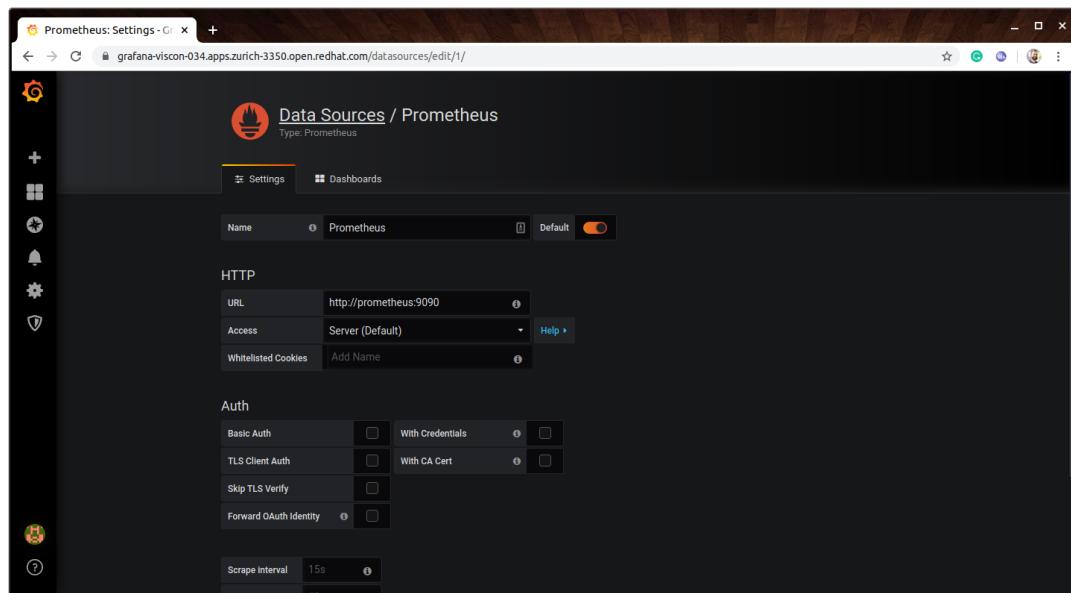
👉 Deploy Grafana

```
oc apply -f strimzi/metrics-demo/grafana.yaml  
oc get routes grafana
```



user: admin, password: admin

✋ Configure the Prometheus Datasource

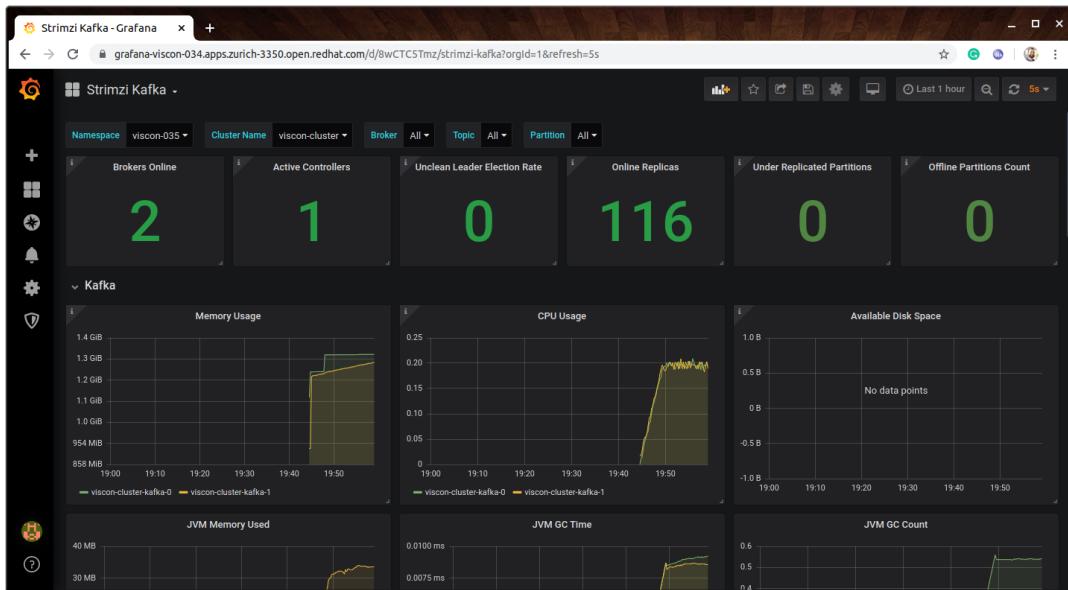


Name: Prometheus, Url: <http://prometheus:9090>



Import the Strimzi Dashboard

Copy `strimzi/metrics-demo/kafka-dashboard.json`



Kafka Eco System

Kafka Connect

Mirror Maker

HTTP Bridge

Managed Kafka

- AWS MSK
- Azure Eventhubs
- Confluent on GCP...

Fireside Chat / Q&A



Contacts and Info

Daniel Streb

Email: daniel.streb@ipt.ch

Linkedin: linkedin.com/in/danistreb

Jonathan Gan

Email: jonathan.gan@ipt.ch

Linkedin: linkedin.com/in/jonganej

IPT Careers Portal

ipt.ch/karriere